

Real Estate House Predictions - Seattle, WA
created by Jaouad



Set environment and libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from scipy import stats
import statsmodels.api as sms
import statsmodels.formula.api as smf
from IPython.display import Image
import seaborn as sns
from sklearn import preprocessing
from sklearn.preprocessing import StandardScaler
%matplotlib inline
```

```
# Read data in dataframe
```

```
df = pd.read_csv("King_County_House_prices_dataset.csv")  
df.head()
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	...	grade
0	7129300520	10/13/2014	221900.0	3	1.00	1180	5650	1.0	NaN	0.0	...	7
1	6414100192	12/9/2014	538000.0	3	2.25	2570	7242	2.0	0.0	0.0	...	7
2	5631500400	2/25/2015	180000.0	2	1.00	770	10000	1.0	0.0	0.0	...	6
3	2487200875	12/9/2014	604000.0	4	3.00	1960	5000	1.0	0.0	0.0	...	7
4	1954400510	2/18/2015	510000.0	3	2.00	1680	8080	1.0	0.0	0.0	...	8

5 rows x 21 columns


```
# Delete columns not needed
```

```
df.drop(["id", "lat", "long", "sqft_basement"], axis=1, inplace=True)
```

```
# Convent date in readable format
```

```
df["date"] = pd.to_datetime(df["date"])  
df["Month"] = df["date"].apply(lambda date: date.month)  
df["Year"] = df["date"].apply(lambda date: date.year)
```

```
df.drop("date", axis=1, inplace =True)
```

```
df.tail(10)
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade	sqft_above
21587	507250.0	3	2.50	2270	5536	2.0	NaN	0.0	3	8	2270
21588	429000.0	3	2.00	1490	1126	3.0	0.0	0.0	3	8	1490
21589	610685.0	4	2.50	2520	6023	2.0	0.0	NaN	3	9	2520
21590	1010000.0	4	3.50	3510	7200	2.0	0.0	0.0	3	9	2600
21591	475000.0	3	2.50	1310	1294	2.0	0.0	0.0	3	8	1180
21592	360000.0	3	2.50	1530	1131	3.0	0.0	0.0	3	8	1530
21593	400000.0	4	2.50	2310	5813	2.0	0.0	0.0	3	8	2310
21594	402101.0	2	0.75	1020	1350	2.0	0.0	0.0	3	7	1020
21595	400000.0	3	2.50	1600	2388	2.0	NaN	0.0	3	8	1600
21596	325000.0	2	0.75	1020	1076	2.0	0.0	0.0	3	7	1020

```
# Show size of the dataset
```

```
df.shape
```

```
(21597, 18)
```

```
# Show overview of the data set
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 21597 entries, 0 to 21596
```

```
Data columns (total 18 columns):
```

#	Column	Non-Null Count	Dtype
0	price	21597 non-null	float64
1	bedrooms	21597 non-null	int64
2	bathrooms	21597 non-null	float64
3	sqft_living	21597 non-null	int64
4	sqft_lot	21597 non-null	int64
5	floors	21597 non-null	float64
6	waterfront	19221 non-null	float64
7	view	21534 non-null	float64
8	condition	21597 non-null	int64
9	grade	21597 non-null	int64
10	sqft_above	21597 non-null	int64
11	yr_built	21597 non-null	int64
12	yr_renovated	17755 non-null	float64
13	zipcode	21597 non-null	int64
14	sqft_living15	21597 non-null	int64
15	sqft_lot15	21597 non-null	int64
16	Month	21597 non-null	int64
17	Year	21597 non-null	int64

```
dtypes: float64(6), int64(12)
```

```
memory usage: 3.0 MB
```

```
# Check if zeros are contained  
df.isnull().values.any()
```

False

```
# View data description and see if there are any outliers  
df.describe()
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view
count	2.159700e+04	21597.000000	21597.000000	21597.000000	2.159700e+04	21597.000000	19221.000000	21534.000000
mean	5.402966e+05	3.373200	2.115826	2080.321850	1.509941e+04	1.494096	0.007596	0.233863
std	3.673681e+05	0.926299	0.768984	918.106125	4.141264e+04	0.539683	0.086825	0.765686
min	7.800000e+04	1.000000	0.500000	370.000000	5.200000e+02	1.000000	0.000000	0.000000
25%	3.220000e+05	3.000000	1.750000	1430.000000	5.040000e+03	1.000000	0.000000	0.000000
50%	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+03	1.500000	0.000000	0.000000
75%	6.450000e+05	4.000000	2.500000	2550.000000	1.068500e+04	2.000000	0.000000	0.000000
max	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+06	3.500000	1.000000	4.000000

```
# Replace the missing data, with the mean value
```

```
df['waterfront'].fillna((df['waterfront'].mean()), inplace=True)  
df['yr_renovated'].fillna((df['yr_renovated'].mean()), inplace=True)  
df['view'].fillna((df['view'].mean()), inplace=True)  
df.dropna(inplace=True)  
df.head()
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade	sqft_above	yr_built
0	221900.0	3	1.00	1180	5650	1.0	0.007596	0.0	3	7	1180	1955
1	538000.0	3	2.25	2570	7242	2.0	0.000000	0.0	3	7	2170	1951
2	180000.0	2	1.00	770	10000	1.0	0.000000	0.0	3	6	770	1933
3	604000.0	4	3.00	1960	5000	1.0	0.000000	0.0	5	7	1050	1965
4	510000.0	3	2.00	1680	8080	1.0	0.000000	0.0	3	8	1680	1987


```
# Convert values to integer with float
```

```
df['waterfront'] = np.round(df['waterfront'])  
df['bathrooms'] = np.round(df['bathrooms'])  
df['yr_renovated'] = np.round(df['yr_renovated'])
```


```
df['floors'] = np.round(df['floors'])  
df['view'] = np.round(df['view'])  
df.head(3)
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade	sqft_above	yr_built
0	221900.0	3	1.0	1180	5650	1.0	0.0	0.0	3	7	1180	1955
1	538000.0	3	2.0	2570	7242	2.0	0.0	0.0	3	7	2170	1951
2	180000.0	2	1.0	770	10000	1.0	0.0	0.0	3	6	770	1933


price




bedrooms




bathrooms




sqft_living




sqft_lot




floors




waterfront




view




condition




grade




sqft_above




yr_built




yr_renovated




zipcode




sqft_living15




sqft_lot15





Month



year







price ~ sqft_living

OLS Regression Results

Dep. Variable:	price		R-squared:	0.493		
Model:	OLS		Adj. R-squared:	0.493		
Method:	Least Squares		F-statistic:	2.097e+04		
Date:	Sun, 06 Jun 2021		Prob (F-statistic):	0.00		
Time:	13:52:08		Log-Likelihood:	-3.0006e+05		
No. Observations:	21597		AIC:	6.001e+05		
Df Residuals:	21595		BIC:	6.001e+05		
Df Model:	1					
Covariance Type:	nonrobust					
	coef	std err	t	P> t 	[0.025	0.975]
Intercept	-4.399e+04	4410.023	-9.975	0.000	-5.26e+04	-3.53e+04
sqft_living	280.8630	1.939	144.819	0.000	277.062	284.664
Omnibus:	14801.942	Durbin-Watson:	1.982			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	542662.604			
Skew:	2.820	Prob(JB):	0.00			
Kurtosis:	26.901	Cond. No.	5.63e+03			


```
# We split our data into training and test data
```

```
X = df[['sqft_living']].values
```

```
y = df['price'].values
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42, test_size=0.25)
```

```
# How good is my prediction
```


```
plt.scatter(X_train, y_train)
```

```
plt.scatter(X_test, y_test, color="red")
```

```
plt.xlabel("sqft_living" )
```

```
plt.ylabel("price")
```

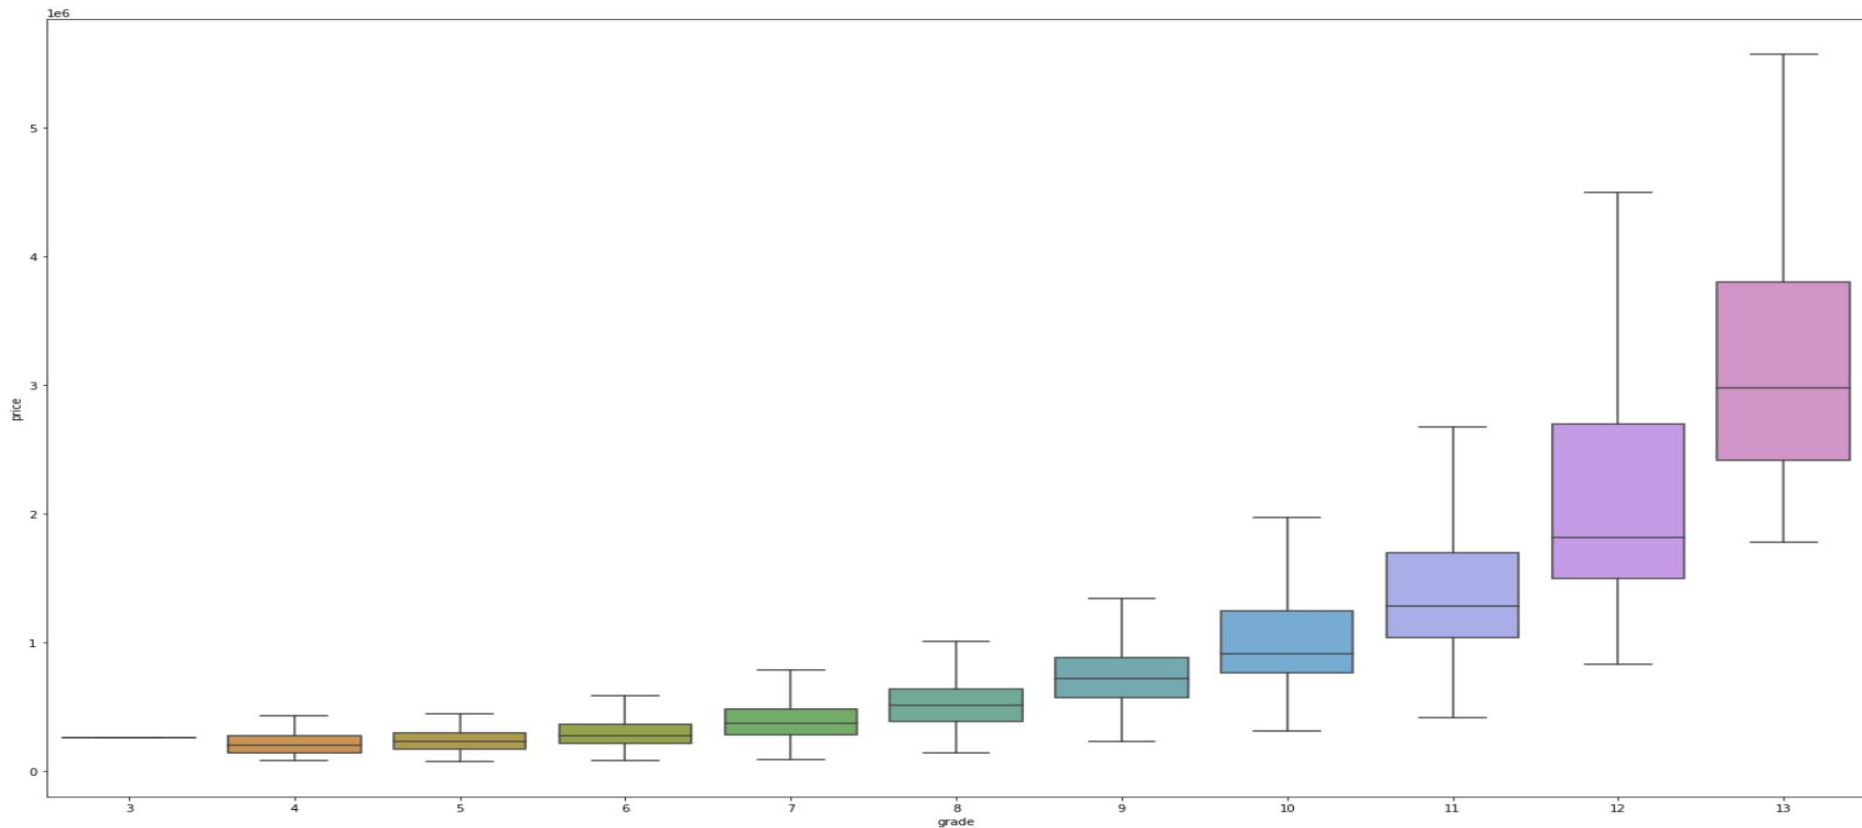
```
plt.show()
```

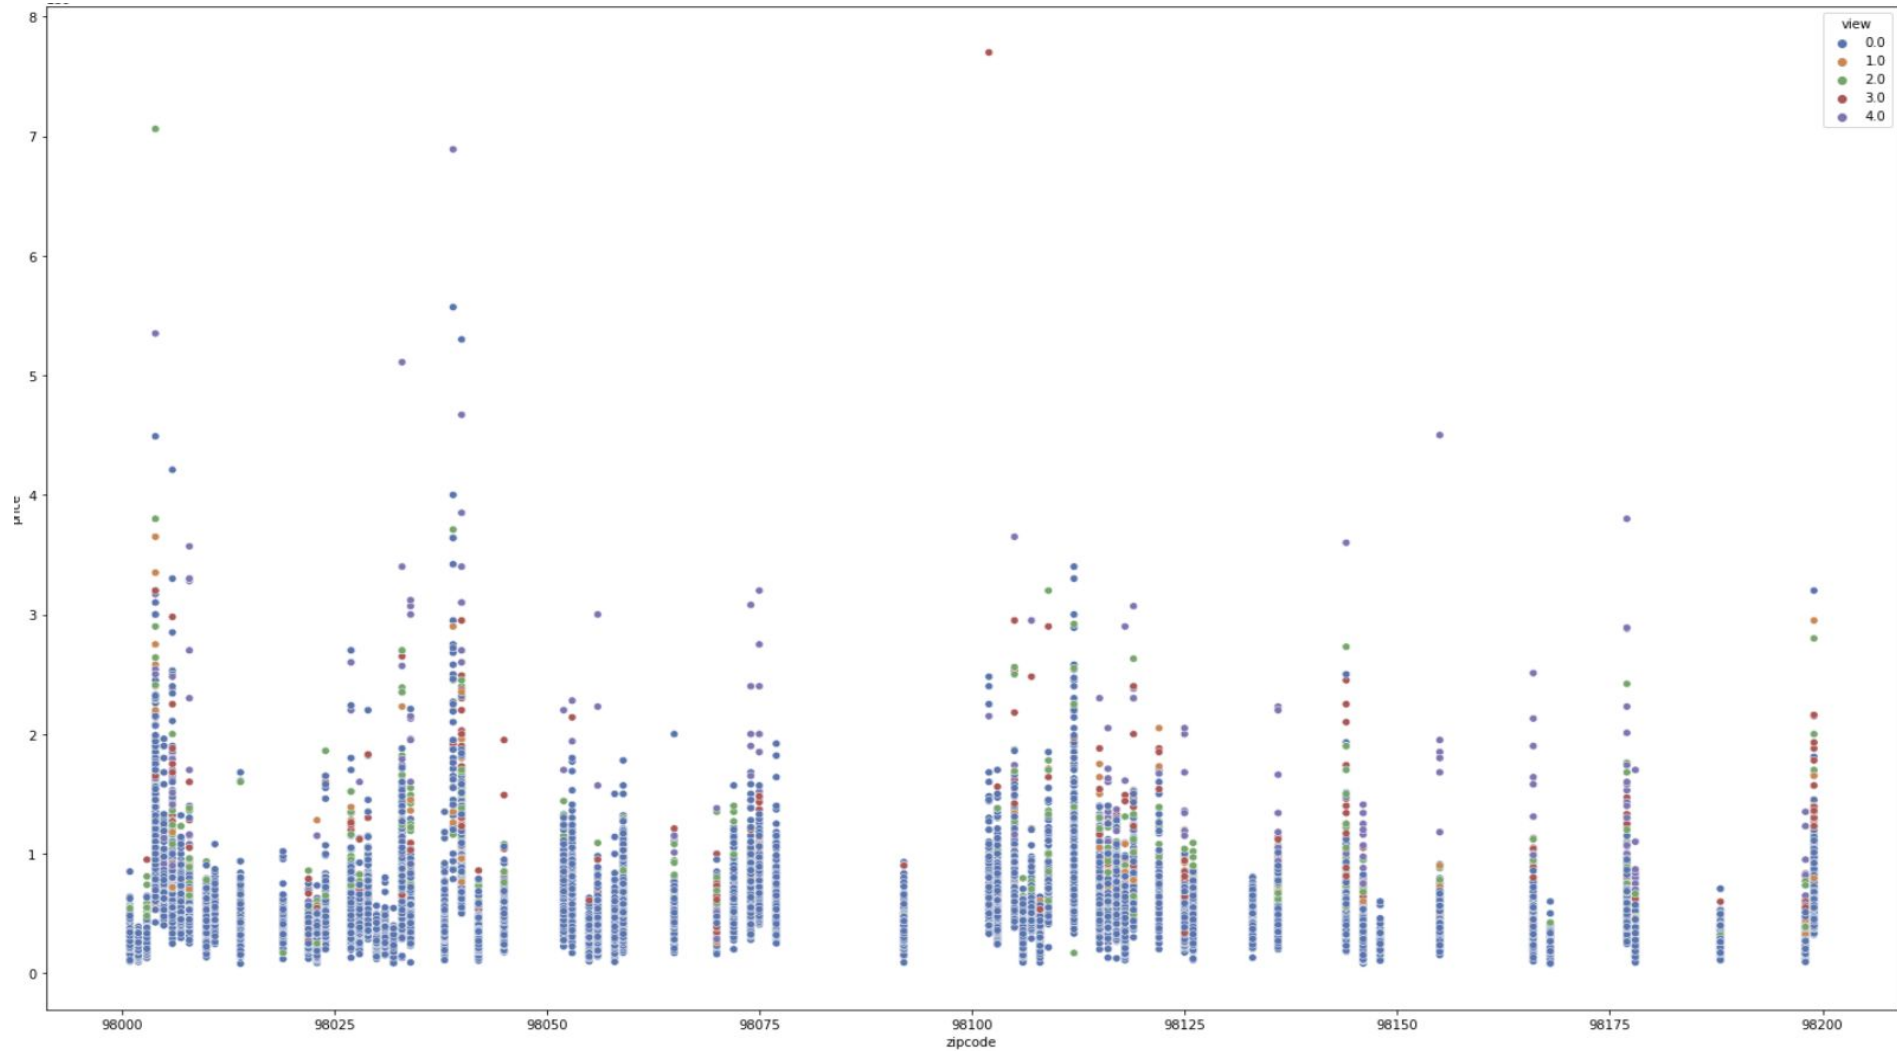


```
# The second strongest correlation between independent and target variable plots
```

```
fig, ax = plt.subplots(figsize=(25,15))  
sns.boxplot(x='grade',y='price',data=df,showfliers=False, ax=ax)
```

```
<AxesSubplot:xlabel='grade', ylabel='price'>
```





```
df = pd.get_dummies(df, columns=["zipcode"])
df.head()
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade	...	zipcode_98146
0	221900.0	3	1.0	1180	5650	1.0	0.0	0.0	3	7	...	0
1	538000.0	3	2.0	2570	7242	2.0	0.0	0.0	3	7	...	0
2	180000.0	2	1.0	770	10000	1.0	0.0	0.0	3	6	...	0
3	604000.0	4	3.0	1960	5000	1.0	0.0	0.0	5	7	...	0
4	510000.0	3	2.0	1680	8080	1.0	0.0	0.0	3	8	...	0


```
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42, test_size=0.25)
```

```
model = LinearRegression()  
model.fit(X_train, y_train)  
#print("Intercept: " +str(model.intercept_))  
#print("Coef: " +str(model.coef_))  
  
print("R2_Score: " + str(model.score(X_test, y_test)))
```

R2_Score: 0.7674287742096019

```
from sklearn.metrics import mean_squared_error  
import math  
print("MSE: " + str(mean_squared_error(y_test, y_test_pred)))  
print("RMSE: " + str(math.sqrt(mean_squared_error(y_test, y_test_pred))))
```

MSE: 31908704848.957275

RMSE: 178630.07823140334

```
df.head(4)
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade	...	zipcode_98146
0	221900.0	3	1.0	1180	5650	1.0	0.0	0.0	3	7	...	0
1	538000.0	3	2.0	2570	7242	2.0	0.0	0.0	3	7	...	0
2	180000.0	2	1.0	770	10000	1.0	0.0	0.0	3	6	...	0
3	604000.0	4	3.0	1960	5000	1.0	0.0	0.0	5	7	...	0

4 rows x 87 columns

```
y_test_pred = model.predict(X_test)  
y_test_pred[0]
```

117539.74891492724

The difference amounts to 104360.00€

```
from sklearn.preprocessing import PolynomialFeatures
```

```
pf = PolynomialFeatures()  
pf.fit(X_train)
```

```
X_train_transformed = pf.transform(X_train)  
X_test_transformed = pf.transform(X_test)
```

```
model = model.fit(X_train_transformed, y_train)  
print(model.score(X_test_transformed, y_test))
```

0.8616663800874034

```
y_test_pred = model.predict(X_test_transformed)  
y_test_pred[0]
```

212651.70591182058

The difference amounts to 9248.00€